PolicyMe: Product Requirements Document (PRD)

Version: 1.0
Date: March 3, 2026
Author: Sirius (PolicyMe Product Team)
Confidentiality: Internal Use Only
Status: Pre-MVP Development Phase

Executive Summary

PolicyMe is an enterprise B2B SaaS platform that transforms unstructured corporate policy documents into a fully navigable knowledge graph powered by GraphRAG (Graph-enhanced Retrieval-Augmented Generation) AI. The platform solves the critical problem of **policy fragmentation** and **compliance risk** by detecting contradictions, answering natural language queries with traceable citations, and providing role-based interfaces for employees, compliance analysts, and system administrators[1][2].

One-Sentence Product Description:
PolicyMe is an AI-powered policy intelligence platform that converts corporate documents into an interactive knowledge graph, enabling instant answers, conflict detection, and audit-ready compliance for enterprise organizations.

Target Market: Mid-to-large enterprises (500+ employees) in regulated industries requiring strict policy compliance (financial services, healthcare, government contractors, manufacturing).

Business Model: Annual subscription SaaS with tiered pricing based on user seats and document volume.

---

1. Problem Statement

1.1 Current State Analysis

Modern enterprises face a systemic failure in corporate policy management:

- **Policy Fragmentation:** Documents scattered across SharePoint, Google Drive, Confluence, and email archives with no centralized access
- **Human Search Inefficiency:** Employees spend 3-5 hours per week searching for policy information, with 40% of searches ending without finding the correct answer
- **Hidden Conflicts:** Contradictory policies exist across HR, IT, Legal, and Compliance departments with no systematic detection mechanism
- **Compliance Blindspots:** Manual policy audits require 40-80 hours per quarter, with high risk of missing critical conflicts before regulatory inspection
- **Knowledge Loss:** When HR/Legal staff leave, institutional knowledge of policy nuances disappears

1.2 Customer Pain Points (by Role)

Employees:

- "I don't know where to find the answer to my policy question"
- "I get different answers from HR vs. my manager"
- "Reading 40-page policy PDFs is painful"

Compliance Analysts:

- "We have no way to detect conflicts until an audit catches them"
- "Manual policy reviews take weeks and still miss edge cases"
- "I can't prove our policies are internally consistent"

IT Administrators:

- "We have policy documents in 6 different systems"
- "There's no access control — everyone sees everything"
- "I have no audit trail of who accessed what policy"

1.3 Quantified Business Impact

- **Productivity Loss:** $8,400 per employee annually (3 hours/week at $60/hour median salary)
- **Compliance Fines:** $2.4M average penalty for policy-related violations in regulated industries
- **Legal Risk:** 67% of employment lawsuits cite policy confusion as contributing factor
- **Onboarding Delays:** New hires take 2-3 weeks to understand policy landscape vs. 2 days with AI assistance

---

2. Product Vision & Strategy

2.1 Vision Statement

"Make every corporate policy instantly accessible, verifiably consistent, and audit-ready through AI-powered knowledge graphs."

PolicyMe will become the single source of truth for enterprise policy intelligence, eliminating fragmentation and enabling proactive compliance management at scale.

2.2 Strategic Positioning

Why PolicyMe Wins:

Dimension	Traditional Solutions	PolicyMe
Search	Keyword-based (Ctrl+F mentality)	GraphRAG with contextual understanding
Conflict Detection	Manual quarterly audits	Real-time AI-powered detection
User Experience	Static PDF repositories	Role-based interactive dashboards
Trust	Black-box AI answers	Full citation transparency with micro-graphs
Integration	Standalone document manager	API-first platform with SSO/HRIS integration

Table 1: Competitive differentiation matrix

2.3 Success Metrics (North Star KPIs)

1. **Primary:** Query Resolution Rate — 85% of employee queries answered without human escalation
2. **Secondary:** Time to Answer — Median query response time under 10 seconds
3. **Tertiary:** Conflict Detection Accuracy — 95% precision on policy contradiction identification
4. **Business:** Customer Retention — 90% annual renewal rate (industry benchmark: 75%)

---

3. Target Users & Personas

3.1 Primary Persona: Sarah (Compliance Analyst)

Demographics:

- Title: Senior Compliance Manager
- Company: 2,000-employee financial services firm
- Age: 38, 10 years experience in compliance

- Reports to: Chief Compliance Officer

Goals:

- Proactively identify policy conflicts before auditors do
- Reduce manual policy review time from 60 hours/quarter to 10 hours
- Generate audit-ready compliance reports with evidence

Pain Points:

- Current process requires reading 200+ policy documents manually
- No systematic way to track policy version changes
- High stress during regulatory audit periods

User Story:

> "As a Compliance Analyst, I need to detect policy conflicts across HR, IT, and Legal documents so that I can remediate issues before our next SOC 2 audit."

3.2 Secondary Persona: Marcus (Employee)

Demographics:

- Title: Mid-level Software Engineer
- Company: 800-employee SaaS startup
- Age: 29, 3 years with company
- Reports to: Engineering Manager

Goals:

- Get instant answers to policy questions (expense limits, PTO rules, remote work)
- Avoid bothering HR for simple questions
- Ensure compliance without reading long documents

Pain Points:

- Current SharePoint has 40 policy PDFs with no search
- HR takes 24-48 hours to respond to Slack questions
- Policies contradict what manager says verbally

User Story:

> "As an Employee, I want to ask natural language questions about company policies so that I can get accurate answers in seconds without waiting for HR."

3.3 Tertiary Persona: David (IT Administrator)

Demographics:

- Title: IT Security Administrator
- Company: 1,500-employee healthcare organization
- Age: 42, 8 years with company
- Reports to: CTO

Goals:

- Centralize policy documents from 5 fragmented systems
- Implement role-based access control (doctors vs. admin staff)
- Maintain audit logs for HIPAA compliance

Pain Points:

- Policies stored in SharePoint, Google Drive, Confluence, and local servers
- No way to control who accesses sensitive HR policies
- Cannot prove access controls during audits

User Story:

> "As an IT Administrator, I need to upload policy documents from multiple sources and assign role-based permissions so that I can maintain HIPAA compliance and audit readiness."

---

4. Core Features & Functional Requirements

4.1 Feature Priority Matrix (MoSCoW Method)

Priority	Features
Must Have (MVP)	Natural language query interface, GraphRAG knowledge graph, Conflict detection engine, Citation micro-graphs, Role-based dashboards (3 roles), Document ingestion (PDF/DOCX), Basic analytics dashboard
Should Have (Post-MVP)	Slack/Teams integration, Advanced filtering (by department/date), Export reports (PDF/CSV), Version control for policies, Real-time collaboration
Could Have (Future)}	Mobile app (iOS/Android), Multi-language support, Advanced NLP training, Browser extension, Workflow automation
Won't Have (Out of Scope)	Document authoring tools, Email client, Custom LLM training, Video content analysis

Table 2: Feature prioritization

4.2 Feature 1: Ask PolicyMe (Natural Language Query Interface)

User Story:
As an Employee, I want to ask questions in plain English so that I can get instant policy answers without reading full documents.

Functional Requirements:

- FR1.1: System shall accept natural language queries via web interface (minimum 500 characters)
- FR1.2: System shall return AI-generated answers within 10 seconds (95th percentile)
- FR1.3: System shall display inline citations with clause references (e.g., "[HR Policy § 3.2]")
- FR1.4: System shall render micro-graph showing query node connected to source document nodes

- FR1.5: System shall allow users to save queries to "My Saved Answers" with tagging
- FR1.6: System shall track query history with timestamps for audit purposes

Acceptance Criteria:

1. Given a query "What is the remote work policy?", when Employee submits, then system returns answer with at least 2 citations within 10 seconds
2. Given an ambiguous query, when system detects multiple interpretations, then system asks clarifying question before answering
3. Given a query with no matching policy, when search fails, then system returns "No policy found" with suggested alternative queries

Technical Specifications:

- Backend: GraphRAG pipeline using Neo4j graph database + OpenAI GPT-4 API[3]
- Vector embeddings: Sentence-BERT for semantic search
- Graph traversal: Multi-hop relationship resolution for context retrieval[3]
- Response format: JSON with answer text, citation IDs, confidence score

4.3 Feature 2: Conflict Detection Engine

User Story:
As a Compliance Analyst, I need to automatically detect policy conflicts so that I can remediate issues before audits.

Functional Requirements:

- FR2.1: System shall analyze uploaded documents and identify contradictory clauses using NLP
- FR2.2: System shall assign conflict severity score (Low/Medium/High) based on impact analysis
- FR2.3: System shall display conflicts in sortable data table with columns: ID, Policy Name, Conflict Type, Severity, Affected

Documents
- FR2.4: System shall highlight conflicting nodes in red on macro-graph visualizer
- FR2.5: System shall generate side-by-side comparison view of conflicting clauses
- FR2.6: System shall allow Analysts to mark conflicts as "Acknowledged" or "Resolved" with notes

Acceptance Criteria:

1. Given two policies with contradictory clauses (e.g., IT bans USB drives but HR issues USB drives to new hires), when documents are uploaded, then system flags conflict with "High" severity
2. Given 50 policy documents, when Analyst runs full scan, then system completes analysis within 5 minutes
3. Given a resolved conflict, when Analyst marks as resolved, then conflict disappears from active table and appears in audit log

Technical Specifications:

- Conflict detection algorithm: Semantic similarity + logical contradiction detection
- NLP pipeline: Entity extraction → Relationship mapping → Contradiction detection
- Graph query: Cypher queries to identify contradictory edges in knowledge graph
- Confidence threshold: Flag conflicts with >70% AI confidence score

4.4 Feature 3: Knowledge Graph Visualizer

User Story:
As a Compliance Analyst, I want to see the entire policy network as an interactive graph so that I can understand relationships between policies.

Functional Requirements:

- FR3.1: System shall render macro-graph with nodes representing policy clauses and edges representing relationships

- FR3.2: System shall color-code nodes (blue = safe, red = conflict, purple = recently updated)
- FR3.3: System shall support zoom/pan interactions with mouse wheel and drag gestures
- FR3.4: System shall display glassmorphic tooltip on node hover with clause metadata
- FR3.5: System shall filter graph by department, document type, or date range
- FR3.6: System shall export graph visualization as PNG/SVG for presentations

Acceptance Criteria:

1. Given 500 policy clauses, when Analyst opens graph visualizer, then system renders all nodes within 2 seconds with 60fps animation
2. Given a conflict node, when Analyst hovers, then tooltip shows conflicting policy names and severity
3. Given a filter selection (e.g., "HR Department only"), when applied, then graph re-renders showing only HR policy nodes

Technical Specifications:

- Visualization library: D3.js or Cytoscape.js[3]
- Graph layout algorithm: Force-directed layout with clustering for departments
- Performance: Canvas rendering for 500+ nodes, WebGL for 1000+ nodes
- Data structure: JSON graph format with nodes and edges arrays

4.5 Feature 4: Role-Based Access Control (RBAC)

User Story:
As an IT Administrator, I need to assign different permissions to users so that employees only see relevant policies.

Functional Requirements:

- FR4.1: System shall support three predefined roles: Employee (read-only), Analyst (read + conflict management), Administrator (full control)

- FR4.2: System shall allow Admins to create custom roles with granular permissions (view, edit, delete, export)
- FR4.3: System shall integrate with corporate SSO (Okta, Azure AD, Google Workspace) for authentication
- FR4.4: System shall route users to role-appropriate dashboards on login
- FR4.5: System shall log all access events (who accessed what document when) for audit trails
- FR4.6: System shall allow Admins to bulk-import users via CSV with role assignments

Acceptance Criteria:

1. Given a new user with Employee role, when user logs in, then system displays Ask PolicyMe interface (not graph visualizer or admin panel)
2. Given an Admin invitation, when new user clicks email link, then system redirects to SSO login and auto-assigns role
3. Given an access log request, when Admin exports audit trail, then CSV includes timestamps, user IDs, document IDs, and actions

Technical Specifications:

- Authentication: OAuth 2.0 with JWT tokens
- Authorization: Role-based middleware checks on every API request
- SSO integrations: SAML 2.0 support for Okta/Azure AD
- Audit logging: Append-only PostgreSQL table with daily backups

4.6 Feature 5: Document Ingestion Pipeline

User Story:
As an IT Administrator, I need to upload policy documents in bulk so that I can migrate from legacy systems.

Functional Requirements:

- FR5.1: System shall accept PDF, DOCX, TXT file formats (max 50MB per file)

- FR5.2: System shall support bulk upload via drag-and-drop interface or ZIP file
- FR5.3: System shall parse documents and extract text with OCR for scanned PDFs
- FR5.4: System shall use NLP to identify entities (policies, clauses, departments) and build knowledge graph
- FR5.5: System shall display progress bar during parsing (estimated time remaining)
- FR5.6: System shall send email notification when processing completes

Acceptance Criteria:

1. Given a 20-page PDF, when Admin uploads, then system completes parsing and graph insertion within 3 minutes
2. Given a scanned PDF with low image quality, when OCR runs, then system achieves 90%+ text extraction accuracy
3. Given 50 documents uploaded simultaneously, when processing begins, then system queues jobs and processes in parallel without crashing

Technical Specifications:

- Text extraction: PyMuPDF for digital PDFs, Tesseract OCR for scanned PDFs
- NLP pipeline: spaCy for entity extraction, OpenAI embeddings for semantic indexing[3]
- Background job processing: Celery task queue with Redis broker
- File storage: AWS S3 with versioning enabled

---

5. Technical Architecture

5.1 System Architecture Diagram

High-Level Components:

- **Frontend:** Next.js 14 (React) with Tailwind CSS, deployed on Vercel
- **Backend API:** Node.js/Express REST API, deployed on AWS ECS (containerized)

- **GraphRAG Engine:** Python microservice with Neo4j graph database + vector embeddings[3]
- **AI/ML Layer:** OpenAI GPT-4 API for query answering, spaCy for NLP
- **Storage:** AWS S3 (documents), PostgreSQL (user data), Neo4j (knowledge graph)
- **Infrastructure:** AWS (compute), Vercel (frontend), Datadog (monitoring)

5.2 Data Flow: Ask PolicyMe Query

1. User submits query via Next.js frontend
2. Frontend sends POST request to `/api/query` endpoint
3. Backend API forwards query to GraphRAG microservice
4. GraphRAG microservice:
    - Generates query embedding using OpenAI embeddings
    - Performs vector similarity search in Neo4j
    - Traverses graph to retrieve related nodes (multi-hop)[3]
    - Constructs context from retrieved clauses
    - Sends context + query to GPT-4 API
    - Receives AI-generated answer with citations
5. GraphRAG returns response to backend API
6. Backend API logs query to audit database
7. Backend sends response to frontend
8. Frontend renders answer, citations, and micro-graph

Latency Targets:

- Total end-to-end: <10 seconds (95th percentile)
- Vector search: <500ms
- Graph traversal: <1 second
- LLM inference: <5 seconds

5.3 Technology Stack

Layer	Technology	Justification
Frontend	Next.js 14, Tailwind CSS, shadcn/ui	Server-side rendering for SEO, Tailwind for rapid UI development
Backend	Node.js, Express, TypeScript	JavaScript ecosystem consistency, fast API development
GraphRAG	Python, Neo4j, LangChain	Neo4j optimized for graph queries, LangChain for RAG pipelines[3]
AI/ML	OpenAI GPT-4, spaCy, Sentence-BERT	Best-in-class LLM, proven NLP library, fast embeddings
Database	PostgreSQL, Neo4j, Redis	Relational for users, graph for policies, Redis for caching
Infrastructure	AWS ECS, S3, RDS, Lambda	Enterprise-grade reliability, auto-scaling, compliance certifications
Monitoring	Datadog, Sentry	Real-time performance metrics, error tracking

Table 3: Technology stack

5.4 Security & Compliance Requirements

- **Authentication:** OAuth 2.0 with SSO (Okta, Azure AD, Google Workspace)
- **Authorization:** Role-based access control (RBAC) with least-privilege principle
- **Data Encryption:** AES-256 encryption at rest (S3, RDS), TLS 1.3 in transit
- **Audit Logging:** Immutable logs for all data access events (HIPAA/SOC 2 compliance)
- **Data Residency:** Region-specific deployments (US, EU, APAC) for GDPR compliance
- **Vulnerability Scanning:** Automated Snyk scans on every code commit

- **Penetration Testing:** Annual third-party security audits

Compliance Certifications (Roadmap):

- Q3 2026: SOC 2 Type I
- Q1 2027: SOC 2 Type II
- Q3 2027: ISO 27001

---

6. User Interface & Design

6.1 Design Philosophy

PolicyMe follows a **Dark Mode Cybersecurity Command Center** aesthetic for the application dashboard, balanced with **Apple-style clarity** for the marketing site[4].

Core Design Principles:

- **Progressive Disclosure:** Show complexity only when needed (Employee sees simple search, Analyst sees full graph)
- **Glassmorphism:** Semi-transparent cards with soft blur effects for modern look
- **Citation Transparency:** Every AI answer includes traceable sources via micro-graphs
- **Single Pane of Glass:** Unified dashboard eliminating context switching

6.2 Key UI Screens

Screen 1: Employee Dashboard (Ask PolicyMe)

- Centered search bar with placeholder: "Ask PolicyMe anything..."
- Recent queries list below search bar
- Right panel: Micro-graph showing citation network
- Top nav: Saved Queries, Settings, Logout

Screen 2: Analyst Dashboard (Conflict Console)

- Left sidebar: Navigation (Dashboard, Graph, Conflicts, Reports)
- Center panel: Conflict data table with sortable columns
- Right panel: Conflict details with side-by-side clause comparison
- Top nav: Export Report, Refresh, Filters

Screen 3: Admin Dashboard (Control Panel)

- Left sidebar: Navigation (Users, Roles, Documents, Audit Logs, Settings)
- Center panel: User management table with Invite button
- Right panel: Recent activity feed
- Top nav: Upload Documents, System Health

6.3 Mobile Considerations

PolicyMe MVP is **desktop-first** (enterprise tool), but supports tablet viewing:

- Desktop ($\geq$1440px): Full layout with expanded sidebar
- Laptop (1024-1439px): Auto-collapsed sidebar, scaled graphs
- Tablet (768-1023px): Sidebar becomes overlay, single-column cards
- Mobile ($<$768px): Not supported (display "Desktop required" message)

---

7. Go-to-Market Strategy

7.1 Target Industries

Primary Targets (Year 1):

1. Financial Services (banks, insurance, fintech) — High compliance requirements, large policy volumes
2. Healthcare (hospitals, pharma) — HIPAA compliance critical, frequent policy updates
3. Government Contractors — Security clearances require strict policy adherence

Secondary Targets (Year 2):

1. Manufacturing — Safety policies, ISO certifications
2. Technology (1000+ employee companies) — Remote work policies, GDPR compliance

7.2 Pricing Strategy

Tier	Starter	Professional	Enterprise
Price/Year	$12,000	$36,000	Custom
User Seats	Up to 100	Up to 500	Unlimited
Document Limit	500 policies	2,000 policies	Unlimited
Conflict Detection	Basic	Advanced	Advanced
Graph Visualizer	Read-only	Interactive	Interactive + Export
SSO Integration	No	Yes	Yes + SAML
Support	Email only	Email + Chat	Dedicated CSM

Table 4: Pricing tiers

Revenue Model: Annual subscription (paid upfront), 20% discount for multi-year contracts.

7.3 Sales Channels

- **Direct Sales:** Outbound SDR team targeting compliance officers at target companies
- **Product-Led Growth:** Free 14-day trial with sample policy dataset
- **Partnerships:** Integrate with HR platforms (BambooHR, Workday) and compliance tools (Drata, Vanta)
- **Content Marketing:** Thought leadership on compliance best practices, SEO for "policy management software"

7.4 Customer Acquisition Cost (CAC) Model

Assumptions:

- Average deal size: $36,000 (Professional tier)
- Sales cycle: 60 days
- Conversion rate: 10% (trial to paid)

CAC Calculation:

- SDR salary: $60,000/year ($5,000/month)
- SDR quota: 5 closed deals/month
- CAC per customer: $5,000 ÷ 5 = $1,000

Target LTV:CAC Ratio: 5:1 (LTV = $36,000 × 3 years × 90% retention = $97,200)

---

8. Development Roadmap & Milestones

8.1 Phase 1: MVP Development (Months 1-4)

Objective: Launch functional MVP with core features for Employee and Analyst roles[5][6].

Key Deliverables:

- Natural language query interface (Ask PolicyMe)
- Basic conflict detection engine
- Knowledge graph visualizer (macro-graph for Analysts, micro-graph for Employees)
- Document ingestion pipeline (PDF/DOCX support)
- Role-based authentication (3 predefined roles)
- Landing page + demo environment

Milestones:

- Month 1: Design system + frontend component library complete
- Month 2: GraphRAG pipeline operational, first query returns answer with citations
- Month 3: Conflict detection algorithm trained, 80% accuracy on test dataset
- Month 4: MVP deployed to staging, beta customer pilot begins

Success Criteria:

- 5 beta customers using platform daily
- 85% query resolution rate
- <10 second query response time

8.2 Phase 2: Feature Expansion (Months 5-8)

Objective: Add enterprise-grade features and integrations based on beta feedback[5][6].

Key Deliverables:

- SSO integration (Okta, Azure AD, Google Workspace)
- Advanced filtering for graph visualizer (by department, date range)
- Export reports (PDF/CSV) with conflict analysis
- Slack/Teams bot for query interface
- Version control for policy documents
- Admin analytics dashboard (query volume, top searched policies)

Milestones:

- Month 5: SSO integration complete, first enterprise pilot begins
- Month 6: Slack bot deployed, 50% of queries shift to Slack channel
- Month 7: Report export feature complete, first compliance audit passed using PolicyMe reports
- Month 8: 20 paying customers, $500K ARR

Success Criteria:

- 90% query resolution rate
- 95% conflict detection accuracy
- 85% customer retention rate

8.3 Phase 3: Scaling & Optimization (Months 9-12)

Objective: Scale infrastructure, optimize performance, achieve SOC 2 compliance[6].

Key Deliverables:

- Auto-scaling infrastructure (AWS ECS with Fargate)
- Performance optimization (sub-5s query latency)
- SOC 2 Type I certification
- Multi-tenancy architecture (data isolation)

- Advanced NLP training (custom entity extraction models)
- Customer success playbook (onboarding, training materials)

Milestones:

- Month 9: Infrastructure supports 50,000 concurrent users
- Month 10: SOC 2 Type I audit begins
- Month 11: Custom NLP models reduce false positive conflicts by 30%
- Month 12: 50 paying customers, $1.5M ARR

Success Criteria:

- 99.9% uptime SLA
- SOC 2 Type I certification achieved
- 90% customer renewal rate

8.4 Phase 4: Enterprise Integration & AI Advancement (Months 13-18)

Objective: Deep enterprise integrations and advanced AI capabilities[6].

Key Deliverables:

- HRIS integrations (BambooHR, Workday) for automatic user provisioning
- Compliance tool integrations (Drata, Vanta) for policy evidence export
- Multi-language support (Spanish, French, German)
- Agentic AI workflows (auto-resolve low-severity conflicts)
- Mobile app (iOS/Android) with simplified Ask PolicyMe interface
- SOC 2 Type II certification

Milestones:

- Month 15: First Fortune 500 customer signs enterprise deal
- Month 16: Multi-language support launches in EU region
- Month 18: 100 paying customers, $5M ARR

Success Criteria:

- 95% query resolution rate
- 98% conflict detection accuracy
- 95% customer retention rate

8.5 Timeline Summary

Phase	Duration	Key Milestones
Phase 1: MVP	Months 1-4	Beta launch, 5 pilot customers, 85% query resolution
Phase 2: Expansion	Months 5-8	20 paying customers, $500K ARR, SSO + integrations
Phase 3: Scaling	Months 9-12	50 customers, $1.5M ARR, SOC 2 Type I
Phase 4: Enterprise	Months 13-18	100 customers, $5M ARR, SOC 2 Type II

Table 5: Development roadmap

---

9. Risk Analysis & Mitigation

9.1 Technical Risks

Risk	Likelihood	Mitigation Strategy
GraphRAG accuracy below 85%	Medium	Extensive NLP training dataset, human-in-the-loop validation for first 1000 queries
Neo4j performance degrades at scale	Low	Benchmark tests with 100K+ nodes, sharding strategy for large customers
OpenAI API cost overruns	Medium	Implement caching layer, switch to self-hosted LLM (Llama 3) for high-volume customers
SSO integration delays	Medium	Prioritize top 3 providers (Okta, Azure AD, Google), use SAML standard library

Table 6: Technical risk matrix

9.2 Business Risks

Risk	Likelihood	Mitigation Strategy
Long enterprise sales cycles (9-12 months)	High	Offer free POC (30 days), focus on mid-market (500-2000 employees) first
Competitor launches similar product	Medium	Focus on citation transparency + conflict detection as differentiators
Customers demand on-premise deployment	Medium	Develop containerized deployment option (Kubernetes) by Month 12
Low user adoption within customer orgs	Medium	Implement Slack bot for frictionless access, gamification (leaderboards)

Table 7: Business risk matrix

9.3 Regulatory Risks

- **GDPR Compliance:** Risk of data residency violations for EU customers
  - Mitigation: Deploy separate EU region (AWS Frankfurt), data localization guarantees in contract
- **AI Liability:** Risk of incorrect policy answers leading to employee violations
  - Mitigation: Clear disclaimer in UI ("AI-generated answer, consult HR for legal advice"), log all queries for audit trail
- **SOC 2 Audit Failure:** Risk of failing first audit due to incomplete controls
  - Mitigation: Hire compliance consultant 3 months before audit, implement Vanta for continuous monitoring

---

10. Success Criteria & KPIs

10.1 Product Success Metrics

Metric	Baseline	Month 6	Month 12
Query Resolution Rate	80%	85%	90%
Median Query Latency	12s	10s	8s
Conflict Detection Accuracy	85%	90%	95%
Daily Active Users (per customer)	N/A	40%	60%
Customer Satisfaction (NPS)	N/A	30	50

Table 8: Product KPIs

10.2 Business Success Metrics

Metric	Baseline	Month 6	Month 12
Paying Customers	0	10	50
Annual Recurring Revenue	$0	$250K	$1.5M
Customer Acquisition Cost	N/A	$2,000	$1,500
Customer Retention Rate	N/A	85%	90%
Average Contract Value	N/A	$25K	$30K

Table 9: Business KPIs

10.3 Leading Indicators

Early signals that predict long-term success:

1. **Time to First Query:** Users ask first question within 5 minutes of onboarding (indicates intuitive UX)
2. **Query Frequency:** Users return to platform 3+ times per week (indicates value creation)
3. **Admin Engagement:** Admins upload 10+ documents in first week (indicates serious evaluation)
4. **Conflict Discovery Rate:** Analysts identify 5+ conflicts in first month (demonstrates ROI)

5. **Stakeholder Expansion:** 3+ departments using platform within 60 days (indicates org-wide adoption)

---

11. Open Questions & Assumptions

11.1 Pending Decisions

1. **Self-Hosted LLM vs. OpenAI API:** Should we build self-hosted option for cost/privacy concerns? (Decision by Month 9)
2. **Light Mode Toggle:** Is dark mode sufficient for MVP, or should we build light mode? (User testing in Month 3)
3. **Mobile App Priority:** Should mobile app be Phase 4 or Phase 5? (Market research in Month 6)
4. **Pricing Model:** Should we offer per-seat pricing or per-document pricing? (A/B test in Month 5)
5. **European Expansion:** When should we deploy EU region for GDPR compliance? (Customer demand assessment in Month 8)

11.2 Critical Assumptions

1. **Assumption 1:** Enterprises are willing to upload sensitive policy documents to cloud platform
   - Validation: SOC 2 certification + data encryption guarantees will satisfy security teams
   - Risk if wrong: Requires pivot to on-premise deployment (6-month delay)
2. **Assumption 2:** Compliance Analysts have authority to purchase software (<$50K budget)
   - Validation: Sales discovery calls with 20 target personas
   - Risk if wrong: Longer sales cycles, need to target C-level buyers
3. **Assumption 3:** GraphRAG provides significant accuracy improvement over basic RAG
   - Validation: Benchmark tests showing 15% accuracy gain[3]
   - Risk if wrong: Simpler architecture could reduce costs
4. **Assumption 4:** Employees will trust AI answers if citations are transparent
   - Validation: User testing with micro-graph prototypes (Month 3)

- Risk if wrong: May need human-in-the-loop validation for all answers

---

12. Appendices

12.1 Competitive Analysis

Feature	PolicyMe	Competitor A (SharePoint)	Competitor B (Guru)
Natural Language Query	Yes (GraphRAG)	No (keyword only)	Yes (basic RAG)
Conflict Detection	Yes (automated)	No	No
Knowledge Graph	Yes (interactive)	No	No
Citation Transparency	Yes (micro-graphs)	N/A	Yes (text links)
Role-Based Dashboards	Yes (3 roles)	Limited	Limited
Pricing	$12K-$36K/year	Included in Office 365	$10/user/month

Table 10: Competitive feature matrix

12.2 Glossary

- **GraphRAG:** Graph-enhanced Retrieval-Augmented Generation, combining knowledge graphs with LLMs for context-aware answers[3]
- **Micro-Graph:** Simplified read-only graph visualization showing 3-5 nodes (for Employee view)
- **Macro-Graph:** Full interactive graph visualization showing entire policy network (for Analyst view)
- **RBAC:** Role-Based Access Control, security model assigning permissions based on user roles

- **SOC 2:** Service Organization Control 2, compliance framework for data security and privacy
- **NPS:** Net Promoter Score, customer satisfaction metric (-100 to +100 scale)
- **ARR:** Annual Recurring Revenue, key SaaS business metric

12.3 References & Research

This PRD was informed by industry best practices in B2B SaaS product development, GraphRAG technical implementation, and enterprise AI roadmaps. Key insights drawn from current standards in policy management software, compliance automation, and knowledge graph architectures[1][2][3][4][5][6].

---

13. Approval & Sign-Off

Stakeholder	Role	Approval Date
Sirius	Product Lead	Pending
[Engineering Lead]	VP Engineering	Pending
[Design Lead]	Head of Design	Pending
[Sales Lead]	VP Sales	Pending
[CEO]	Chief Executive Officer	Pending

Table 11: Approval matrix

Document Status: Draft — Pending stakeholder review
Next Review Date: March 10, 2026
Version History: 1.0 (Initial draft)

References

[1] Atlassian. (2026). How to create a product requirements document (PRD). https://www.atlassian.com/agile/product-management/requirements

[2] Infrasity. (2025). The ultimate PRD guide for your B2B SaaS success in 2025. https://www.infrasity.com/blog/b2b-saas-prd

[3] Memgraph. (2024). Prerequisites for GraphRAG systems. https://memgraph.com/docs/ai-ecosystem/graph-rag/prerequisites

[4] Codelevate. (2025). The ultimate PRD guide for your B2B SaaS success in 2025. https://www.codelevate.com/blog/the-ultimate-prd-guide-for-your-b2b-saas-success-in-2025

[5] Neontri. (2026). Enterprise AI roadmap 2026: Implementation framework. https://neontri.com/blog/enterprise-ai-roadmap/

[6] Promethium AI. (2025). Enterprise AI implementation roadmap and timeline. https://promethium.ai/guides/enterprise-ai-implementation-roadmap-timeline/